

PERFORMANCE TESTING

Date	21 May 2023
Team ID	NM2023TMID17607
Project Name	Cancer Mortality and Incidence rates classification using ML

Cancer Death Rates

CANCER DEATH RATE	METRICS																																										
Model	CONFUSION MATRIX	ACCURACY SCORE	CLASSIFICATION REPORT																																								
Decision Tree	<pre>print(confusion_matrix(test_y,y_pred)) [[20 0 0 0] [0 119 0 4] [0 0 0 1] [2 6 0 399]]</pre>	<pre>print(accuracy_score(test_y,y_pred)) 0.9764065335753176</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>Data too sparse to predict a trend</td><td>0.91</td><td>1.00</td><td>0.95</td><td>20</td></tr><tr><td>falling</td><td>0.95</td><td>0.97</td><td>0.96</td><td>123</td></tr><tr><td>rising</td><td>0.00</td><td>0.00</td><td>0.00</td><td>1</td></tr><tr><td>stable</td><td>0.99</td><td>0.98</td><td>0.98</td><td>407</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.98</td><td>551</td></tr><tr><td>macro avg</td><td>0.71</td><td>0.74</td><td>0.72</td><td>551</td></tr><tr><td>weighted avg</td><td>0.98</td><td>0.98</td><td>0.98</td><td>551</td></tr></table>		precision	recall	f1-score	support	Data too sparse to predict a trend	0.91	1.00	0.95	20	falling	0.95	0.97	0.96	123	rising	0.00	0.00	0.00	1	stable	0.99	0.98	0.98	407	accuracy			0.98	551	macro avg	0.71	0.74	0.72	551	weighted avg	0.98	0.98	0.98	551
	precision	recall	f1-score	support																																							
Data too sparse to predict a trend	0.91	1.00	0.95	20																																							
falling	0.95	0.97	0.96	123																																							
rising	0.00	0.00	0.00	1																																							
stable	0.99	0.98	0.98	407																																							
accuracy			0.98	551																																							
macro avg	0.71	0.74	0.72	551																																							
weighted avg	0.98	0.98	0.98	551																																							
Random Forest	<pre>print(confusion_matrix(test_y,y_pred)) [[20 0 0 0] [0 122 0 1] [0 0 0 1] [0 8 0 399]]</pre>	<pre>print(accuracy_score(test_y,y_pred)) 0.9818511796733213</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>Data too sparse to predict a trend</td><td>1.00</td><td>1.00</td><td>1.00</td><td>20</td></tr><tr><td>falling</td><td>0.94</td><td>0.99</td><td>0.96</td><td>123</td></tr><tr><td>rising</td><td>0.00</td><td>0.00</td><td>0.00</td><td>1</td></tr><tr><td>stable</td><td>1.00</td><td>0.98</td><td>0.99</td><td>407</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.98</td><td>551</td></tr><tr><td>macro avg</td><td>0.73</td><td>0.74</td><td>0.74</td><td>551</td></tr><tr><td>weighted avg</td><td>0.98</td><td>0.98</td><td>0.98</td><td>551</td></tr></table>		precision	recall	f1-score	support	Data too sparse to predict a trend	1.00	1.00	1.00	20	falling	0.94	0.99	0.96	123	rising	0.00	0.00	0.00	1	stable	1.00	0.98	0.99	407	accuracy			0.98	551	macro avg	0.73	0.74	0.74	551	weighted avg	0.98	0.98	0.98	551
	precision	recall	f1-score	support																																							
Data too sparse to predict a trend	1.00	1.00	1.00	20																																							
falling	0.94	0.99	0.96	123																																							
rising	0.00	0.00	0.00	1																																							
stable	1.00	0.98	0.99	407																																							
accuracy			0.98	551																																							
macro avg	0.73	0.74	0.74	551																																							
weighted avg	0.98	0.98	0.98	551																																							
K NEAREST NEIGHBORS	<pre>print(confusion_matrix(test_y,y_pred)) [[5 2 0 13] [0 42 0 81] [0 0 0 1] [8 27 0 372]]</pre>	<pre>print(accuracy_score(test_y,y_pred)) 0.7604355716878403</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>Data too sparse to predict a trend</td><td>0.38</td><td>0.25</td><td>0.30</td><td>20</td></tr><tr><td>falling</td><td>0.59</td><td>0.34</td><td>0.43</td><td>123</td></tr><tr><td>rising</td><td>0.00</td><td>0.00</td><td>0.00</td><td>1</td></tr><tr><td>stable</td><td>0.80</td><td>0.91</td><td>0.85</td><td>407</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.76</td><td>551</td></tr><tr><td>macro avg</td><td>0.44</td><td>0.38</td><td>0.40</td><td>551</td></tr><tr><td>weighted avg</td><td>0.73</td><td>0.76</td><td>0.74</td><td>551</td></tr></table>		precision	recall	f1-score	support	Data too sparse to predict a trend	0.38	0.25	0.30	20	falling	0.59	0.34	0.43	123	rising	0.00	0.00	0.00	1	stable	0.80	0.91	0.85	407	accuracy			0.76	551	macro avg	0.44	0.38	0.40	551	weighted avg	0.73	0.76	0.74	551
	precision	recall	f1-score	support																																							
Data too sparse to predict a trend	0.38	0.25	0.30	20																																							
falling	0.59	0.34	0.43	123																																							
rising	0.00	0.00	0.00	1																																							
stable	0.80	0.91	0.85	407																																							
accuracy			0.76	551																																							
macro avg	0.44	0.38	0.40	551																																							
weighted avg	0.73	0.76	0.74	551																																							
Naïve Bayes	<pre>print(confusion_matrix(test_y,y_pred)) [[19 0 0 1] [0 107 0 16] [0 0 0 1] [40 35 9 323]]</pre>	<pre>print(accuracy_score(test_y,y_pred)) 0.8148820326678766</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>Data too sparse to predict a trend</td><td>0.32</td><td>0.95</td><td>0.48</td><td>20</td></tr><tr><td>falling</td><td>0.75</td><td>0.87</td><td>0.81</td><td>123</td></tr><tr><td>rising</td><td>0.00</td><td>0.00</td><td>0.00</td><td>1</td></tr><tr><td>stable</td><td>0.95</td><td>0.79</td><td>0.86</td><td>407</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.81</td><td>551</td></tr><tr><td>macro avg</td><td>0.51</td><td>0.65</td><td>0.54</td><td>551</td></tr><tr><td>weighted avg</td><td>0.88</td><td>0.81</td><td>0.84</td><td>551</td></tr></table>		precision	recall	f1-score	support	Data too sparse to predict a trend	0.32	0.95	0.48	20	falling	0.75	0.87	0.81	123	rising	0.00	0.00	0.00	1	stable	0.95	0.79	0.86	407	accuracy			0.81	551	macro avg	0.51	0.65	0.54	551	weighted avg	0.88	0.81	0.84	551
	precision	recall	f1-score	support																																							
Data too sparse to predict a trend	0.32	0.95	0.48	20																																							
falling	0.75	0.87	0.81	123																																							
rising	0.00	0.00	0.00	1																																							
stable	0.95	0.79	0.86	407																																							
accuracy			0.81	551																																							
macro avg	0.51	0.65	0.54	551																																							
weighted avg	0.88	0.81	0.84	551																																							

Support
Vector
Classification

```
print(confusion_matrix(test_y,y_pred))
```

```
[[ 6  0  0 14]
 [ 1 79  0 43]
 [ 0  0  0  1]
 [ 6 14  2 385]]
```

```
print(accuracy_score(test_y,y_pred))
```

0.852994555353902

```
print(classification_report(test_y,y_pred))
```

	precision	recall	f1-score	support
Data too sparse to predict a trend	0.46	0.30	0.36	20
falling	0.85	0.64	0.73	123
rising	0.00	0.00	0.00	1
stable	0.87	0.95	0.91	407
accuracy			0.85	551
macro avg	0.55	0.47	0.50	551
weighted avg	0.85	0.85	0.85	551

CANCER DEATH RATE	HYPERPARAMETER TUNING
MODELS	GridSearchCV
Decision Tree	<pre>#Decision Tree params = {'max_leaf_nodes': list(range(2, 100)), 'min_samples_split': [2, 3, 4]} grid_search_cv = GridSearchCV(DecisionTreeClassifier(random_state=42), params, verbose=1, cv=3) grid_search_cv.fit(train_x, train_y) Fitting 3 folds for each of 294 candidates, totalling 882 fits GridSearchCV(cv=3, estimator=DecisionTreeClassifier(random_state=42), param_grid={'max_leaf_nodes': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, ...], 'min_samples_split': [2, 3, 4]}, verbose=1) score_dt=grid_search_cv.best_score_ param_dt=grid_search_cv.best_params_ print('Best Score of Decision Tree:',score_dt) print('Best Parameters of Decision Tree:',param_dt) Best Score of Decision Tree: 0.9812462189957653 Best Parameters of Decision Tree: {'max_leaf_nodes': 19, 'min_samples_split': 4}</pre>
Random Forest	<pre>#Random Forest clf=GridSearchCV(RandomForestClassifier(),{'n_estimators':[1,5,10]},cv=5,return_train_score=False) clf.fit(x,y) score_rf=clf.best_score_ param_rf=clf.best_params_ print('Best Score of Random Forest:',score_rf) print('Best parameters of Random Forest:',param_rf) Best Score of Random Forest: 0.9346052360338074 Best parameters of Random Forest: {'n_estimators': 10}</pre>
K Nearest Neighbors	<pre>#KNN k_range = list(range(1, 31)) param_grid = dict(n_neighbors=k_range) # defining parameter range grid = GridSearchCV(KNeighborsClassifier(), param_grid, cv=10, scoring='accuracy', return_train_score=False, verbose=1) # fitting the model for grid search grid_search=grid.fit(train_x,train_y) score_knn=grid_search.best_score_ param_knn=grid_search.best_params_ print('Best Score of KNN:',score_knn) print('Best parameters of KNN:',param_knn) Fitting 10 folds for each of 30 candidates, totalling 300 fits Best Score of KNN: 0.7513581599123768 Best parameters of KNN: {'n_neighbors': 12}</pre>
Naïve Bayes	<pre>#NaiveBayes from sklearn.model_selection import RepeatedStratifiedKFold cv_method = RepeatedStratifiedKFold(n_splits=5, n_repeats=3, random_state=999) from sklearn.preprocessing import PowerTransformer params_NB = {'var_smoothing': np.logspace(0,-9, num=100)} gs_NB = GridSearchCV(GaussianNB(), param_grid=params_NB, cv=cv_method,verbose=1,scoring='accuracy') Data_transformed = PowerTransformer().fit_transform(x) gs_NB.fit(Data_transformed, y) Fitting 15 folds for each of 100 candidates, totalling 1500 fits GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=3, n_splits=5, random_state=999), estimator=GaussianNB(), param_grid={'var_smoothing': array([1.00000000e+00, 8.11130831e-01, 6.57933225e-01, 5.33669923e-01, 4.32876128e-01, 3.51119173e-01, 2.84803587e-01, 2.31012970e-01, 1.87381742e-01, 1.51991108e-01, 1.23284674e-01, 1.00000000e-01, 8.11130831e-02, 6.57933225e-02, 5..., 1.23284674e-07, 1.00000000e-07, 8.11130831e-08, 6.57933225e-08, 5.33669923e-08, 4.32876128e-08, 3.51119173e-08, 2.84803587e-08, 2.31012970e-08, 1.87381742e-08, 1.51991108e-08, 1.23284674e-08, 1.00000000e-08, 8.11130831e-09, 6.57933225e-09, 5.33669923e-09, 4.32876128e-09, 3.51119173e-09, 2.84803587e-09, 2.31012970e-09, 1.87381742e-09, 1.51991108e-09, 1.23284674e-09, 1.00000000e-09])}), scoring='accuracy', verbose=1) score_NB=gs_NB.best_score_ param_NB=gs_NB.best_params_ print('Best Score of Naive Bayes:',score_NB) print('Best parameters of Naive Bayes:',param_NB) Best Score of Naive Bayes: 0.9218054696626126 Best parameters of Naive Bayes: {'var_smoothing': 0.0001232846739442066}</pre>

Scores_comparison

	Models	Scores before hyperparameter tuning	Scores after hyperparameter tuning
0	Decision Tree	0.976407	0.981246
1	Random Forest	0.981851	0.934605
2	KNN	0.760436	0.751358
3	Naive Bayes	0.814882	0.921805

Models_comparison

	Models	Scores	Parameters
0	Decision Tree	0.981246	{'max_leaf_nodes': 19, 'min_samples_split': 4}
1	Random Forest	0.934605	{'n_estimators': 10}
2	KNN	0.751358	{'n_neighbors': 12}
3	Naive Bayes	0.921805	{'var_smoothing': 0.0001232846739442066}

Cancer Incidence Rates

CANCER INCIDENCE RATES	METRICS		
MODELS	CONFUSION MATRIX	ACCURACY SCORE	CLASSIFICATION REPORT
Decision Tree	<pre>print(confusion_matrix(test_y,y_pred))</pre> <pre>[[10 0 0 0] [0 26 0 0] [0 0 5 0] [4 0 0 476]]</pre>	<pre>print(accuracy_score(test_y,y_pred))</pre> <pre>0.9923224568138196</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <pre> precision recall f1-score support</pre> <pre>confidential 0.71 1.00 0.83 10</pre> <pre>falling 1.00 1.00 1.00 26</pre> <pre>rising 1.00 1.00 1.00 5</pre> <pre>stable 1.00 0.99 1.00 480</pre> <pre>accuracy 0.93 1.00 0.99 521</pre> <pre>macro avg 0.99 0.99 0.99 521</pre> <pre>weighted avg</pre>
Random Forest	<pre>print(confusion_matrix(test_y,y_pred))</pre> <pre>[[10 0 0 0] [0 26 0 0] [0 0 5 0] [0 0 0 480]]</pre>	<pre>print(accuracy_score(test_y,y_pred))</pre> <pre>1.0</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <pre> precision recall f1-score support</pre> <pre>confidential 1.00 1.00 1.00 10</pre> <pre>falling 1.00 1.00 1.00 26</pre> <pre>rising 1.00 1.00 1.00 5</pre> <pre>stable 1.00 1.00 1.00 480</pre> <pre>accuracy 1.00 1.00 1.00 521</pre> <pre>macro avg 1.00 1.00 1.00 521</pre> <pre>weighted avg</pre>
K Nearest Neighbors	<pre>print(confusion_matrix(test_y,y_pred))</pre> <pre>[[0 0 0 10] [1 1 0 24] [0 0 0 5] [0 3 0 477]]</pre>	<pre>print(accuracy_score(test_y,y_pred))</pre> <pre>0.9174664107485605</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <pre> precision recall f1-score support</pre> <pre>confidential 0.00 0.00 0.00 10</pre> <pre>falling 0.25 0.04 0.07 26</pre> <pre>rising 0.00 0.00 0.00 5</pre> <pre>stable 0.92 0.99 0.96 480</pre> <pre>accuracy 0.29 0.26 0.26 521</pre> <pre>macro avg 0.86 0.92 0.89 521</pre> <pre>weighted avg</pre>
Naïve Bayes	<pre>print(confusion_matrix(test_y,y_pred))</pre> <pre>[[10 0 0 0] [0 21 0 5] [0 0 2 3] [1 4 2 473]]</pre>	<pre>print(accuracy_score(test_y,y_pred))</pre> <pre>0.9712092130518234</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <pre> precision recall f1-score support</pre> <pre>confidential 0.91 1.00 0.95 10</pre> <pre>falling 0.84 0.81 0.82 26</pre> <pre>rising 0.50 0.40 0.44 5</pre> <pre>stable 0.98 0.99 0.98 480</pre> <pre>accuracy 0.81 0.80 0.80 521</pre> <pre>macro avg 0.97 0.97 0.97 521</pre> <pre>weighted avg</pre>
Support Vector Classification	<pre>print(confusion_matrix(test_y,y_pred))</pre> <pre>[[2 0 0 8] [0 18 0 8] [0 0 3 2] [0 0 1 479]]</pre>	<pre>print(accuracy_score(test_y,y_pred))</pre> <pre>0.963531669865643</pre>	<pre>print(classification_report(test_y,y_pred))</pre> <pre> precision recall f1-score support</pre> <pre>confidential 1.00 0.20 0.33 10</pre> <pre>falling 1.00 0.69 0.82 26</pre> <pre>rising 0.75 0.60 0.67 5</pre> <pre>stable 0.96 1.00 0.98 480</pre> <pre>accuracy 0.93 0.62 0.70 521</pre> <pre>macro avg 0.96 0.96 0.96 521</pre> <pre>weighted avg</pre>

CANCER INCIDENCE RATES	HYPERPARAMETER TUNING
MODELS	GridSearchCV
Decision Tree	<pre>#Decision Tree params = {'max_leaf_nodes': list(range(2, 100)), 'min_samples_split': [2, 3, 4]} grid_search_cv = GridSearchCV(DecisionTreeClassifier(random_state=42), params, verbose=1, cv=3) grid_search_cv.fit(train_x, train_y) Fitting 3 folds for each of 294 candidates, totalling 882 fits GridSearchCV(cv=3, estimator=DecisionTreeClassifier(random_state=42), param_grid={'max_leaf_nodes': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, ...], 'min_samples_split': [2, 3, 4]}, verbose=1) score_dt=grid_search_cv.best_score_ param_dt=grid_search_cv.best_params_ print('Best Score of Decision Tree:',score_dt) print('Best Parameters of Decision Tree:',param_dt) Best Score of Decision Tree: 0.992948717948718 Best Parameters of Decision Tree: {'max_leaf_nodes': 9, 'min_samples_split': 2}</pre>
Random Forest	<pre>#Random Forest clf=GridSearchCV(RandomForestClassifier(),{'n_estimators':[1,5,10]},cv=5,return_train_score=False) clf.fit(x,y) score_rf=clf.best_score_ param_rf=clf.best_params_ print('Best Score of Random Forest:',score_rf) print('Best parameters of Random Forest:',param_rf) Best Score of Random Forest: 0.8013927319682714 Best parameters of Random Forest: {'n_estimators': 1}</pre>
K Nearest Neighbors	<pre>#KNN k_range = list(range(1, 31)) param_grid = dict(n_neighbors=k_range) # defining parameter range grid = GridSearchCV(KNeighborsClassifier(), param_grid, cv=10, scoring='accuracy', return_train_score=False,verbose=1) # fitting the model for grid search grid_search=grid.fit(train_x,train_y) score_knn=grid_search.best_score_ param_knn=grid_search.best_params_ print('Best Score of KNN:',score_knn) print('Best parameters of KNN:',param_knn) Fitting 10 folds for each of 30 candidates, totalling 300 fits Best Score of KNN: 0.9243580743580744 Best parameters of KNN: {'n_neighbors': 11}</pre>
Naïve Bayes	<pre>#NaiveBayes from sklearn.model_selection import RepeatedStratifiedKFold cv_method = RepeatedStratifiedKFold(n_splits=5, n_repeats=3, random_state=999) from sklearn.preprocessing import PowerTransformer params_NB = {'var_smoothing': np.logspace(0,-9, num=100)} gs_NB = GridSearchCV(GaussianNB(), param_grid=params_NB, cv=cv_method,verbose=1,scoring='accuracy') Data_transformed = PowerTransformer().fit_transform(x) gs_NB.fit(Data_transformed, y) Fitting 15 folds for each of 100 candidates, totalling 1500 fits GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=3, n_splits=5, random_state=999), estimator=GaussianNB(), param_grid={'var_smoothing': array([1.00000000e+00, 8.11130831e-01, 6.57933225e-01, 5.33669923e-01, 4.32876128e-01, 3.51119173e-01, 2.84803587e-01, 2.31012970e-01, 1.87381742e-01, 1.51991108e-01, 1.23284674e-01, 1.00000000e-01, 8.11130831e-02, 6.57933225e-02, 5... 1.23284674e-07, 1.00000000e-07, 8.11130831e-08, 6.57933225e-08, 5.33669923e-08, 4.32876128e-08, 3.51119173e-08, 2.84803587e-08, 2.31012970e-08, 1.87381742e-08, 1.51991108e-08, 1.23284674e-08, 1.00000000e-08, 8.11130831e-09, 6.57933225e-09, 5.33669923e-09, 4.32876128e-09, 3.51119173e-09, 2.84803587e-09, 2.31012970e-09, 1.87381742e-09, 1.51991108e-09, 1.23284674e-09, 1.00000000e-09])]), scoring='accuracy', verbose=1) score_NB=gs_NB.best_score_ param_NB=gs_NB.best_params_ print('Best Score of Naive Bayes:',score_NB) print('Best parameters of Naive Bayes:',param_NB) Best Score of Naive Bayes: 0.9769326846215336 Best parameters of Naive Bayes: {'var_smoothing': 0.0003511191734215131}</pre>
Support Vector Classification	<pre>#SVC clf=GridSearchCV(SVC(),{'C':[1,5,10],'kernel':['rbf','linear']],cv=5,return_train_score=False) clf.fit(x,y) score_svc=clf.best_score_ param_svc=clf.best_params_ print('Best Score of SVC:',score_svc) print('Best parameters of SVC:',param_svc) Best Score of SVC: 0.9490534495480538 Best parameters of SVC: {'C': 10, 'kernel': 'linear'}</pre>

	Models	Scores before hyperparameter tuning	Scores after hyperparameter tuning
0	Decision Tree	0.992322	0.992949
1	Random Forest	0.996161	0.801393
2	KNN	0.917466	0.924359
3	Naive Bayes	0.971209	0.976933
4	SVC	0.963532	0.949053

Models_comparison

	Models	Scores	Parameters
0	Decision Tree	0.992949	{'max_leaf_nodes': 9, 'min_samples_split': 2}
1	Random Forest	0.801393	{'n_estimators': 1}
2	KNN	0.924359	{'n_neighbors': 11}
3	Naive Bayes	0.976933	{'var_smoothing': 0.0003511191734215131}
4	SVC	0.949053	{'C': 10, 'kernel': 'linear'}